


Chapter 5: Week 3, Session A — Introduction to Microsoft Copilot & Prompting Essentials



"The quality of your question determines the quality of your thinking."

— Peter Drucker

There is a version of Microsoft Copilot that is, frankly, underwhelming.

You type: *"Write me a report on commercial real estate."*

Copilot dutifully returns four paragraphs of polished, completely generic prose that could have been written about any bank in any city in any year. It reads like a Wikipedia summary. You could have found better with a three-second web search. You close the window, return to your keyboard, and quietly conclude that AI is probably overhyped.

There is a second version of Microsoft Copilot — the version that the top performers in every organization that has deployed it eventually discover. In this version, you type a different kind of prompt. Copilot pulls from your actual emails, your recent meeting transcripts, the spreadsheet your analyst shared yesterday, and

the memo your compliance officer sent last Friday. It gives you a first draft that reflects your tone, incorporates your institution's specific numbers, and accounts for the regulatory context you operate in. It saves you forty-five minutes. Then it saves you another forty-five. Then it rewrites the way you work.

The difference between those two outcomes is not the software. It is the prompt.

This chapter is about the architecture that makes the second outcome possible — how Microsoft 365 Copilot actually works, where to find it, how it connects to your organizational data — and then the five prompting techniques that separate average users from power users. Master these, and Copilot stops being a curiosity and starts being your most capable colleague.

1. The Copilot Landscape: Not All Copilots Are Equal


Before we discuss prompting, we need to clear up a confusion that trips up most new users:

Microsoft has multiple products with

"Copilot" in the name, and they are not the same thing.

This is not a minor technical footnote. If you understand the distinction, you will know exactly which tool to reach for and why. If you don't, you'll end up frustrated that "Copilot" doesn't have access to your files — because you're using the wrong one.

FIGURE 4.2: Microsoft 365 Copilot Integration in Commercial Banking – Cross-Platform Workflow



Professional Banking Education Series: Utilizing AI for Efficient Loan Analysis

Think of it this way: imagine a hospital. They have a public information desk at the front entrance — anyone can walk in, ask questions, and get general health information. That's useful. But when you need a specialist who actually knows your medical history, who can look at your chart and make recommendations

based on *your* specific situation, you see a physician with access to your full record. That is a fundamentally different kind of help.

Copilot Chat (free, included in your Microsoft account) is the public information desk. It is web-connected, it can answer general questions, it can help you write a cover letter or explain a news story — but by default, it does not have access to your BankUnited emails, your SharePoint files, your Teams conversations, or any other organizational data. It is general-purpose. It is genuinely useful. But it is not your specialist.

Microsoft 365 Copilot (the paid license add-on that BankUnited has deployed) is the physician with your chart. It knows your organizational context. It connects to Microsoft Graph — the intelligence layer that maps your emails, your meetings, your chats, and your documents — and it uses that context to give you responses grounded in *your* actual work. When you ask it to summarize the key points from your conversation with a relationship manager last

Tuesday, it can actually do that. When you ask it to draft a memo based on the spreadsheet your analyst shared this morning, it can pull that file.

Copilot Agents are a third category: specialized, purpose-built AI assistants configured to handle specific workflows. A Copilot Agent might be designed specifically to handle loan document review, or to automate a compliance reporting workflow, or to answer customer questions from a specific product knowledge base. These are more advanced and, for most BankUnited users in this course, are the horizon toward which we're building — the payoff of foundational mastery.

:

{tab-item} Copilot Chat (Free)

What it is: A general-purpose AI assistant connected to the web.

What it can do:

- Answer general questions using web search
- Help draft documents, emails, and summaries

- Explain concepts, analyze text you paste in
- Work with content you explicitly share in the conversation

What it cannot do:

- Access your BankUnited emails or calendars
- Pull from your Teams chats or SharePoint files
- Reference organizational documents without you pasting them in

Best for: General research, personal productivity, drafting when you have content to share manually.

Access: copilot.microsoft.com

What it is: An AI assistant grounded in your organizational context via Microsoft Graph.

What it can do:

- *Access your emails, meetings, files, and chats (only what you have permission to see)*

- *Draft documents referencing actual BankUnited files in SharePoint/OneDrive*
- *Summarize Teams meetings you attended*
- *Create Excel analyses from data in your tables*
- *Search across your entire M365 ecosystem*

What it cannot do:

- *Access data you don't have permission to view*
- *See other employees' private files or emails*
- *Go outside your Microsoft 365 security boundary*

Best for: *Everything in this course. This is your primary tool.*

Access: *m365.cloud.microsoft*

What they are: Purpose-built AI assistants configured for specific workflows or knowledge bases.

What they can do:

- *Automate specific, repeatable workflows*
- *Answer questions from a defined knowledge base (e.g., BankUnited product library)*
- *Handle multi-step processes without step-by-step prompting*
- *Integrate with specific systems beyond M365*

What they require:

- *Configuration by IT or a designated Copilot admin*
- *Clear definition of scope and data sources*

Best for: Advanced use cases — loan processing assistance, compliance Q&A, customer service automation.

Access: Through Microsoft Copilot Studio
(admin-configured)

....

The Security Point That Should Give You Confidence

One of the most common concerns we hear from BankUnited colleagues is: "If Copilot can see all my organizational data, can it share my clients' information with someone else?"

*The answer is no — and the architecture makes this structurally impossible. Microsoft 365 Copilot only surfaces information that the **signed-in user already has permission to access**. If a loan officer doesn't have permission to view another relationship manager's client files, Copilot cannot show them those files. The same access controls, conditional access policies, and compliance frameworks that govern your M365 environment govern Copilot. Your customer*

data stays within the Microsoft 365 service boundary. It does not leave your tenant, it is not used to train Microsoft's AI models, and it is not accessible to other organizations.


This is not marketing language. It is the architecture. Understanding it removes a legitimate concern and lets you focus on what Copilot can actually do.

2. The Front Door: m365.cloud.microsoft

Every powerful tool has an entry point. For Microsoft 365 Copilot, that entry point is **m365.cloud.microsoft** — the unified Microsoft 365 Copilot portal that brings together everything in one place.

Think of m365.cloud.microsoft the way you think of a bank's main branch lobby: it's not the only way to access services (you can also use the ATM, the drive-through, the mobile app), but it's

the hub — the place where everything connects and where you can do the most, all in one location.



What you'll find at m365.cloud.microsoft:

The portal gives you access to the full Microsoft 365 Copilot experience from your browser — the same capability available through the Microsoft 365 Copilot desktop app and the mobile app. You can start conversations, manage ongoing threads, use the Prompt Gallery (more on this shortly), and access all Copilot integrations from a single authenticated session.

Three ways to access M365 Copilot:

- **Web:** m365.cloud.microsoft — works in any modern browser, no installation required

- **Desktop:** The Microsoft 365 Copilot app, available through your Microsoft 365 installation
- **Mobile:** The Microsoft 365 Copilot mobile app (iOS and Android) — the same Copilot, on your phone, with the same access to your organizational data

Bookmark it now. Open *m365.cloud.microsoft* in your browser and add it to your bookmarks bar. You'll use it daily. It takes four seconds, and it removes one more friction point between you and a tool that should feel like second nature.

Copilot embedded in the apps you already use:

m365.cloud.microsoft is the standalone experience. But Microsoft 365 Copilot also lives inside every app in your M365 suite — and for most BankUnited workflows, you'll use it there rather than switching to the portal.

- **In Word:** Copilot appears in the document margin. Ask it to draft a section, expand a

paragraph, rewrite in a different tone, or summarize the document. You can reference a specific file from SharePoint by typing `/` and the file name in your prompt.

- **In Excel:** Copilot appears in the ribbon on the Home tab. It can analyze data in a table, generate formulas, create charts, identify trends, and surface outliers — without you needing to write a single formula manually.
- **In Outlook:** Copilot can draft email replies, summarize long email threads, and flag action items from your inbox.
- **In Teams:** Copilot can summarize meetings you missed, recap decisions made during a call, and list action items — pulling from the meeting transcript in real time.
- **In PowerPoint:** Copilot can generate slide decks from a document, reorganize presentations, and suggest design improvements.

The pattern across all of these is consistent: **you stay in the tool you're already using, and Copilot shows up as a natural part of the workflow.** You don't need to stop what you're doing and go somewhere else.

A Word About File Access in Word and Copilot

*Microsoft's official guidance is important here: for Copilot to access a file when you're prompting in Word, **that file must be stored in SharePoint or OneDrive** — not just on your local hard drive. If a file lives only on your desktop or a local folder, Copilot cannot reach it. This is by design — it's part of the security boundary.*

The practical implication for BankUnited: store your working documents in SharePoint or OneDrive, not local drives. This is good practice regardless of AI — it ensures documents are backed up, version-controlled, and accessible to your team — but it's essential for Copilot to work as designed.

Microsoft 365 Copilot Search:

The portal also includes **Microsoft 365 Copilot Search** — a universal search capability that works across all your M365 apps and connected

third-party data. Instead of searching separately in Outlook, then SharePoint, then Teams, Copilot Search finds relevant content across your entire organizational ecosystem in a single query. For a relationship manager trying to find everything related to a specific client across emails, call notes, shared documents, and team chats — this alone is a significant capability.

3. Work IQ — The Intelligence Layer That Grounds Copilot in Your Context

Here is the feature that makes Microsoft 365 Copilot genuinely different from a general-purpose AI chatbot — and it is important enough that it deserves its own section.

When you ask a general AI tool a question, it answers from what it knows from training data — the internet, books, articles. That's valuable. But it doesn't know *you*. It doesn't know what happened in your portfolio review meeting last

Tuesday. It doesn't know the memo your compliance team circulated last week. It doesn't know the deal terms your relationship manager just shared via email.

Microsoft 365 Copilot uses **Microsoft Graph** to close that gap.

Microsoft Graph is the intelligence layer that maps your organizational context. Think of it as the connective tissue of your Microsoft 365 environment — it knows which emails you've sent and received, which meetings you've attended, which documents you've created and accessed, which chats you've had in Teams, and which files your colleagues have shared with you. It maps the relationships between all of this information and maintains that map in real time.

When you submit a prompt to Microsoft 365 Copilot, the system does something called **grounding**: before generating a response, Copilot accesses Microsoft Graph within your tenant to pull in relevant context from your actual work. If you ask "Summarize what I need to prepare for the Riverside Holdings meeting tomorrow," Copilot doesn't just give you generic meeting prep advice — it looks at the meeting

invitation, pulls the agenda, finds relevant emails from recent weeks, locates documents shared with the attendees, and synthesizes all of that into a personalized briefing.

This process — the prompt going in, the grounding against your Microsoft Graph data, and the response coming out — all happens within your organization's Microsoft 365 service boundary. Your data does not leave your tenant.

The Permission Principle: How Grounding Stays Safe

*The grounding process is governed by a principle that is worth understanding clearly: **Copilot only surfaces information the signed-in user already has permission to access.***

This is not a policy statement. It is structural. Copilot doesn't have a special administrative view of your organization's data that bypasses your existing access controls. It uses Microsoft Graph with the same permissions as your account. If you don't have access to a file in SharePoint,

Copilot cannot include that file in your responses. If a colleague's calendar is set to private, Copilot cannot see the details of their appointments.

For BankUnited, this means: the existing security model that governs your M365 environment — including Conditional Access, MFA, and all other compliance controls — governs Copilot's access. No new exposure. The same guardrails, applied to AI-assisted work.

What grounding means in practice at BankUnited:

- A commercial lender asks Copilot to draft a credit memo. Copilot pulls the relevant loan application documents from SharePoint, the financial statements shared by the client via email, and the notes from the last underwriting committee discussion — all automatically, because they're already in the user's M365 environment.
- An operations manager asks Copilot to flag any compliance-related items from the last 30

days of email. Copilot searches the user's inbox for relevant threads and surfaces a prioritized summary.

- A relationship manager asks Copilot to prepare talking points for a client renewal conversation. Copilot finds the account history, recent correspondence, and product documentation in SharePoint and synthesizes key points for the meeting.

None of these require the user to manually attach files or paste content into the prompt. The intelligence layer does the retrieval automatically — because it already knows your work context.

The analogy that makes this click:

Imagine you have a brilliant new analyst who has been working alongside you for three months. They've attended every meeting with you, read every email you've sent and received, reviewed every document in your shared drive. When you ask them to help you prepare for a client meeting, they don't need you to explain the history of the relationship — they already know it. They synthesize what they know into what you need.


That is what Microsoft Graph grounding does for Copilot. You don't need to brief it every time. It already knows your context. You just need to ask the right question.

The honest caveat: the analogy breaks down when it comes to judgment and creativity. Your analyst builds genuine understanding and relationship intuition over time. Copilot's "knowledge" is a structured map of your data — incredibly useful, but not a substitute for the contextual judgment you bring. The output it generates is always a first draft, not a final verdict. Microsoft itself is explicit about this in its product documentation: *"Remember that Copilot generates a draft. You'll need to verify and modify details to make sure it's accurate and fits your tone and style."*

Read that sentence carefully. It is not a liability disclaimer buried in fine print. It is honest product guidance from Microsoft — and it is good professional practice.

4. The Microsoft 365 Copilot Prompt Gallery and Skilling Center

Before we get to the five prompting techniques, you need to know about two resources Microsoft has built specifically to help you get better at this faster.



The Copilot Prompt Gallery lives at m365.cloud.microsoft/copilot-prompts (<https://m365.cloud.microsoft/copilot-prompts>).

It is a curated library of ready-to-use prompts organized by job function (finance, HR, operations, sales, legal) and by application (Word, Excel, Outlook, Teams). Every prompt in the gallery has been tested and verified — these

are not examples someone invented in a slide deck. They are prompts that produce results in the actual M365 Copilot environment.

For BankUnited users, the gallery is your starting point, not your ceiling. Browse the finance and operations categories to find proven prompts you can adapt to your specific workflows. Then modify them using the five techniques in the next section to make them sharper, more specific, and more powerful.

The Microsoft 365 Copilot Skilling Center is the official learning hub at adoption.microsoft.com/copilot/skilling-center (<https://adoption.microsoft.com/copilot/skilling-center/>). It includes structured learning paths, scenario-based guides, and role-specific content — all officially maintained by Microsoft. If you ever want to go deeper than this course takes you, the Skilling Center is the verified source.

How to Use the Prompt Gallery Right Now

- 1. Go to m365.cloud.microsoft/copilot-prompts*
- 2. Filter by your primary job function*
- 3. Find three prompts that address tasks you do regularly*
- 4. Try each one this week in the relevant M365 app*
- 5. Note what worked, what didn't, and what you'd adjust*


This is a fifteen-minute investment that will pay dividends in your first week of real Copilot use.

5. The Five Prompting Techniques

Now we get to the core of this chapter.

Microsoft 365 Copilot has access to your organizational context. It has sophisticated language models coordinating its responses. It is, architecturally speaking, a remarkable tool. But none of that matters if you don't know how to talk to it.

Prompting is not a technical skill. It is a communication skill. The same way a well-constructed question to a colleague gets you a better answer than a vague one, a well-constructed prompt gets you better output from Copilot. The techniques below are not tricks — they are frameworks for clarity. Learn them, practice them, and they become second nature.




Technique 1: Role-Based Prompting

The core idea: Before asking Copilot to do something, tell it who it is.

This sounds almost comically simple. It is also the technique that produces the most immediate and dramatic improvement in output quality for most users.

Here is why it works: language models like the one powering Microsoft 365 Copilot don't have a single fixed "voice" or perspective. They adapt their reasoning, vocabulary, tone, and framing based on the context they're given. When you assign a role — *"Act as a veteran commercial real estate underwriter with 20 years of experience"* — you are not just changing the tone. You are shifting the entire conceptual frame from which Copilot approaches the task. The criteria it applies, the risks it looks for, the language it uses, the depth of analysis it attempts — all of these shift to match the assigned role.

Role-Based Prompting — Give Copilot a Persona



The analogy: Think about what happens when you call a law firm and ask a question. If you call the main line and speak to a receptionist, you get a polite general answer. If you ask to speak to the partner who handles commercial real estate transactions, you get a different kind of answer entirely — one that draws on specific expertise, that sees the risks you didn't think to ask about, that speaks in the language of your actual problem.

Role-Based Prompting is how you route your request to the right expert, even when the expert is an AI.

Where the analogy breaks down: a real expert has genuine experiential judgment, client relationship memory, and liability stake in their

advice. Copilot's role adoption is sophisticated pattern matching, not lived expertise. You still bring the judgment. Copilot brings the synthesis and first draft.

The revolution this enables: Consider what it means for a compliance analyst to prompt Copilot as a *"skeptical bank examiner reviewing this policy document for OCC alignment gaps."* Or for a relationship manager to prompt as a *"CFO of a mid-market manufacturing company who is skeptical about borrowing right now."* Or for an operations manager to prompt as a *"Six Sigma process improvement consultant looking for waste and bottlenecks."* Each role unlocks a different analytical lens — applied instantly, at scale, to your actual BankUnited materials.

Prompting template:

Act **as** a [specific role **with** relevant background].
Your task **is to** [specific task].
[Additional context **about** what you need.]

BankUnited examples:

Act **as** a cynical, veteran CRE underwriter **with** 20 years **of** experience **at** a Southeast regional bank. Review **the** following **property** financial summary **and** flag **every** assumption you would push **back on in** an underwriting committee discussion.

[Paste **or reference the** financial summary]

Act **as a** BSA/AML compliance officer preparing **for a** quarterly regulatory review. Read **the** following internal **process** document **and** identify **any** language that could raise questions **from an** OCC examiner. Be specific about which examiner guidelines **each** concern maps **to**.

Act **as** a relationship manager who has known this client **for** five years. Based **on the** following account summary **and** recent correspondence, draft a renewal conversation agenda **that** acknowledges **the** relationship history **and** anticipates **the** client's likely concerns **about** rate.

Role-Based Prompting Power Move


Add a behavioral instruction alongside the role: "Be direct. Do not soften your feedback. Flag problems explicitly."

Copilot's default tendency is to be somewhat diplomatic. In professional banking contexts, you often need it to be blunt. The role assignment plus the behavioral instruction together produce output that reads like a tough internal review, not a polished PR document.

Technique 2: Chain-of-Thought Reasoning

The core idea: Ask Copilot to show its work before giving you the final answer.

This is one of the most counterintuitive techniques, because the instinct is always to ask for the answer — not the reasoning process. But in complex analytical tasks, asking for reasoning first produces dramatically better final answers. Here is why.



When you ask Copilot directly for a conclusion, it pattern-matches toward the most statistically likely answer given the context. When you ask it to reason through the problem step-by-step, it builds each conclusion on the previous one — and that structured process tends to produce more coherent, internally consistent analysis. More importantly, it makes the reasoning visible, which means you can catch errors in the logic before they propagate into the conclusion.

The analogy: Think about a credit memo versus a credit decision. A credit decision (approve/decline, with a rate) is the conclusion. A credit memo is the reasoning chain — market analysis, borrower financials, collateral assessment, risk mitigants, stress testing — that makes the decision reviewable, auditable, and

defensible. When Copilot shows its reasoning chain, you get the memo instead of just the verdict. That is almost always more useful, because you can engage with the analysis rather than just accepting or rejecting the conclusion.

Where the analogy breaks down: a credit memo is built on real due diligence by humans who are accountable for their analysis. Copilot's reasoning chain is sophisticated generation — it should be reviewed for logical consistency, not treated as due diligence.

The prompting template:

Walk me **through** your reasoning step-**by**-step **before** giving me **the** final answer.
[Then state **the** task clearly.]

BankUnited examples:

I need **to** evaluate whether **to** recommend our bank's SBA 504 program **to** a prospect who currently has a conventional commercial term loan **with** us coming up **for** renewal.

Walk **me** through your reasoning **step-by-step**: what factors should I consider, what are the trade-offs, **and** what questions **do** I need answered before making a recommendation? **Then** give **me** your final recommendation framework.

A commercial deposit client who generates \ \$4.2M **in** average deposits **is** asking us **to** match a competitor's offer that includes a higher yield, dedicated treasury services, **and** a dedicated relationship manager. We currently provide standard treasury services.

Walk **me** through your reasoning **on** whether **to** negotiate, match, **or let** this client leave – **then** give **me** a prioritized recommendation **with** the **key** risks clearly stated.

At the model architecture level, Chain-of-Thought prompting works because it forces the language model to generate tokens that represent intermediate reasoning before generating the final answer token sequence. Each step in the chain conditions the next step — creating a more constrained, coherent generation path than a direct-to-conclusion prompt allows.


In plain language: when you ask for the answer directly, Copilot takes a shortcut. When you ask for the reasoning first, it has to build the bridge piece by piece, and those intermediate pieces keep it on track. The final answer that comes out of a reasoning chain is usually more defensible than one that appeared directly.

This is also why Chain-of-Thought is particularly valuable in compliance-adjacent banking tasks: the reasoning chain is auditable. You can show it to a senior colleague or a regulator and explain why the analysis landed where it did.

Technique 3: Few-Shot Prompting

The core idea: Show Copilot what "good" looks like before asking it to produce something new.

This technique is borrowed directly from machine learning — "few-shot learning" is the ability of a model to generalize from a small number of examples. Applied to prompting, it means you give Copilot two or three strong examples of the output you want, and then ask it to produce a new output in the same style.



The analogy: Think about how you onboard a new junior analyst who has strong technical skills but no institutional knowledge. You don't hand them a style guide and tell them to write

like the senior team. You hand them three or four examples of excellent credit memos written by your best people and say: *"Write at this level. Match this structure. Aim for this depth."* They look at the examples, internalize the pattern, and produce work that fits the institutional standard.

This is exactly what few-shot prompting does. The examples are your institutional standard, communicated directly to Copilot.

Where the analogy breaks down: your analyst internalizes examples through genuine understanding, not pattern matching. Copilot extracts statistical patterns from your examples — which means if your examples have a systematic flaw or bias, Copilot will replicate that flaw. The "garbage in, garbage out" principle applies here more directly than in most other techniques.

The prompting template:

Here are [number] examples of [type of output].

[Example 1 – labeled "Strong" or "Weak" as appropriate]

[Example 2]

[Example 3]

Now [state the task]. Match the style, depth, and structure of the strong examples.

BankUnited examples:

Here are three examples of client relationship summary memos that our team considers high quality, and one example of a weaker memo.

[EXAMPLE 1 – Strong: paste text]

[EXAMPLE 2 – Strong: paste text]

[EXAMPLE 3 – Weak: paste text]

Now write a relationship summary memo for the following client using the structure and depth of the strong examples.

Here is the raw account data:

[paste account information]

Here are two examples of the way our compliance team prefers to write policy exceptions – concise, specific, with clear risk mitigation language and no ambiguous qualifiers.

[EXAMPLE 1]

[EXAMPLE 2]

Now write a policy exception memo for the following situation:

[describe situation]

Few-Shot Prompting and Confidential Data

When using examples to train Copilot's style, be thoughtful about what you paste in. Real client names, account numbers, and non-public financial information should be anonymized in your examples before pasting them into a Copilot prompt. Use placeholder names (Client A, Company X) and rounded figures in your examples. The style instruction works just as well with anonymized content, and it keeps your prompting practice in alignment with BankUnited's data handling standards.

Advanced application — the "Four Examples" method:


For complex output like proposals or memos, try four examples: two strong, one acceptable, one weak. Ask Copilot to explain what makes the strong examples better before generating new output. This forces a brief analysis step (a variation of Chain-of-Thought) that tends to sharpen the generation significantly.

Technique 4: Reverse Prompting

The core idea: Instead of you asking Copilot a question, ask Copilot to ask *you* questions.

This is the technique that most surprises people when they first try it — because it inverts the natural instinct. We assume we know what we need and we try to describe it. The problem is that we often don't know what we need precisely enough to describe it well. The resulting prompt is vague, the output is generic, and we conclude that Copilot isn't that useful.

Reverse Prompting short-circuits this by letting Copilot do the requirement-gathering.



The analogy: Think about the difference between walking into a tailor with a vague request ("I need a suit for an important meeting") versus answering a tailor's structured questions ("What is the occasion? What is the dress code of the venue? What is your typical posture preference for jacket length? Do you prefer a two-button or three-button closure?"). The tailor who asks questions produces a significantly better outcome than the tailor who takes your vague description and starts cutting fabric. Reverse Prompting makes Copilot the skilled tailor.

Where the analogy breaks down: a skilled tailor's questions are based on deep craft knowledge and an understanding of *your specific situation*. Copilot's questions are good but may not always surface the most important uncertainty for your specific task. After it asks its questions, you should also add anything it didn't think to ask.

The prompting template:

I need your help **with** [general description **of** task].

Before you **start**, ask me **any** questions you need **to** perfectly understand my request **and** produce exactly what I need. Don't **proceed until I've** answered.

BankUnited examples:

I need **to write a** presentation **for** our quarterly business review **with the** commercial lending leadership team. Before you **start**, ask me **any** questions you need **to** understand exactly what this presentation should accomplish, who will be **in the** room, what data I have available, **and** what **format** works best **for** our team.

I want **to** draft a proposal letter **to** a prospect we've been courting for six months in the commercial construction space. Before writing anything, ask me everything you need to know about the relationship history, the prospect's **current** banking situation, what we're proposing, and the tone I want.

The Reverse Prompting Power Move

After Copilot asks its questions and you answer them, add one more line: "Is there anything else you need to know, or any assumption you're making that I should verify?" This second-level check often surfaces a critical variable that the first round of questions missed — particularly on complex, multi-stakeholder tasks.

When Reverse Prompting works best:

- Complex, multi-part deliverables (proposals, presentations, strategy memos)


- Tasks where the audience or context matters significantly to the output
 - Situations where you're not sure exactly what you want — but you know it when you see it
 - Tasks you're doing for the first time with Copilot and haven't yet developed a strong prompt template for
-

Technique 5: The Sparring Partner

The core idea: Ask Copilot to push back on your ideas, not agree with them.

This is the most advanced technique in the set — and for many banking professionals, the most valuable. The default behavior of any AI assistant is to be helpful, which in practice means it tends to be agreeable. It will draft your memo, it will refine your proposal, it will polish your arguments. What it will not do, unless you specifically ask it to, is tell you where your thinking is weak.

The Sparring Partner technique flips this. You give Copilot a role defined by skepticism, and then you present your best thinking and ask it to attack it.



The analogy: Think about moot court in law school. Law students argue their cases against experienced lawyers who deliberately try to find every hole in the argument. The goal is not to win the moot court — it is to discover every weakness in the case before the real court appearance, when the stakes are real. The sparring session is where the case gets stronger.

For banking professionals, the Sparring Partner technique is how you pressure-test your thinking before a credit committee review, a client negotiation, a board presentation, or a regulatory examination. Copilot, assigned the role of your toughest critic, will find the holes that your own confirmation bias is inclined to skip over.

Where the analogy breaks down: a real sparring partner in moot court knows the law, has genuine judgment, and cares about the outcome. Copilot's pushback is sophisticated, but it is still pattern-matched critique rather than genuine adversarial reasoning. It may miss highly specific or contextual weaknesses that a real expert would catch. Use it as a first filter, not the only one.

The prompting template:

Play **the** role **of** [specific skeptical counterpart].
I am going **to** present [my idea/proposal/analysis].
Push **back** hard. Don't flatter **me**. Identify **every** weakness, **every** assumption, **every** place **where** my reasoning **is** vulnerable. Be direct.

[Present your idea/proposal/analysis]

BankUnited examples:

Play **the** role **of** a skeptical CFO **of** a \ \$200M revenue manufacturing company who **is** cautious **about** taking **on** additional debt **in the** current rate environment. I'm going **to** pitch you **on** refinancing your current term loan **with** BankUnited **and** adding a revolving credit facility. Push **back** hard. Don't accept **my** assumptions. Tell **me** what objections you'd actually raise.

Here **is my** pitch: [present your pitch]

Play **the** role **of** a lead OCC examiner conducting a safety **and** soundness review. I am going **to** present our commercial **real** estate concentration management framework. Identify **every** area **where the** examiner **is** likely **to** ask hard questions, **where** our documentation **is** thin, **and where** our reasoning may **not** hold up under scrutiny.

Here **is the** framework: [paste **the** framework]

Play the role of a credit committee member who has seen five similar credits go bad in the last two years. **I'm** presenting a \ \$7.5M construction loan for a mixed-use development in Broward County. **Find** the weaknesses in my credit story. **Do** not soften your feedback.


Here is my credit summary: [paste the summary]

The Sparring Partner and Overconfidence

The Sparring Partner technique is extraordinarily useful — but it carries one risk worth naming. Copilot's pushback is sophisticated enough that it may feel comprehensive. It may not be. There are scenarios, contextual nuances, and relationship factors that a human expert will surface that Copilot will miss. Run your ideas through the Sparring Partner to improve them. Then also run them through an actual human colleague you trust. The combination — AI critique followed by human review — produces the best outcome.

6. The Goldilocks Zone of Prompt Length

Now that you understand the five techniques, there is one meta-principle that governs all of them: **the length and precision of your prompt determines the quality of your output, and there is a precision sweet spot — not too short, not too long.**



The "too short" failure mode:

"Write a memo about our Q2 commercial lending performance."

This prompt has no context (who is reading it? what do they already know?), no specific goal (is this for leadership? compliance? the board?), no format guidance (length? tone? structure?), and no constraints (what should it include? what should it avoid?). Copilot will generate something. It will be polished. It will be almost

entirely useless for your specific purpose, because it is optimized for the generic version of the task rather than your version.

The "too long" failure mode:

Some users, having learned that detail helps, overcompensate. They write prompts that are five paragraphs of background, instruction, counter-instruction, caveats, and competing requirements. At a certain level of complexity, this actually degrades output — the model gets pulled in multiple directions by conflicting guidance, and the result is incoherent or over-hedged.

The Goldilocks zone — the four-component prompt:

The research and practitioner experience on prompting converges on a consistent structure for most business tasks. A good prompt has four components:

1. **Context:** Who you are, what situation you're in, what Copilot needs to know to frame its response appropriately
2. **Goal:** What you want Copilot to produce —

specific, not vague

3. **Format:** How you want the output structured (bullet list, formal memo, executive summary, table, numbered recommendations, etc.)
4. **Constraints:** What to include, what to avoid, tone, length, any requirements that bound the output

These four components can usually be covered in three to five sentences, or a short paragraph. That is the sweet spot.

The BankUnited four-component prompt in action:

Context: I'm a commercial relationship manager at BankUnited preparing **for** a renewal conversation **with** a \ \$15M commercial real estate client who has been **with** the bank **for** eight years. Their portfolio has performed well, but a competitor has approached them **with** a more aggressive rate.

Goal: Draft talking points **for** the renewal conversation that emphasize our relationship history **and** service quality without immediately conceding **on** rate.

Format: A bulleted list **of** five **to** seven talking points, **each with** a one-sentence expansion.

Constraints: Keep the tone confident but **not** defensive. **Do not** include any specific rate numbers **or** commitments — I'll handle those in the room. The client **is** sophisticated **and** has heard every standard pitch before.

That prompt takes thirty seconds to write. The output it produces is immediately usable — not a generic draft that requires fifteen minutes of editing to make relevant.

```
flowchart LR
    A["Context\n(Who + Situation)"] --> E["The Precision\nSweet Spot"]
    B["Goal\n(Specific Output)"] --> E
    C["Format\n(Structure)"] --> E
    D["Constraints\n(Boundaries)"] --> E
    E --> F["Professional-Grade\nCopilot Output"]
    style E fill:#1a73e8,color:#fff
    style F fill:#f4a400,color:#fff
```

A note on iterating:

You rarely need to get the perfect prompt on the first try. Copilot conversations are threaded — your follow-up messages have context from everything that came before. Think of prompting as a conversation: start with a good four-component prompt, review the output, and then refine with specific follow-up instructions.

"That's good — now make the tone more direct"
or *"Remove the second and fourth points and expand the third into two separate bullets"* or
"Rewrite this for a non-banker who doesn't know our internal terminology." The follow-up prompt is often where the real refinement happens.

7. Try This: Run the Same Banking Question Five Ways

Here is where the theory becomes practice. This exercise produces one of the clearest demonstrations of how much the technique matters — more than the question itself.

The baseline question:

"How should BankUnited approach a commercial real estate client who is approaching loan maturity during a period of elevated rates and declining collateral values?"

This is a real strategic question that any commercial lender or relationship manager at BankUnited might face. It is specific enough to be meaningful, general enough to work without confidential client data.

Run it five times — once with each technique:

:

{tab-item} Technique 1: Role-Based

Act **as a** veteran workout officer **at a** regional bank **with** 15 years **of** experience managing distressed commercial real estate credits.
How should BankUnited approach **a** commercial real estate client who is approaching loan maturity during **a** period **of** elevated rates **and** declining collateral values?

What to notice: The response should have a notably more specific, risk-focused, and action-oriented character than a generic answer. It should reference specific workout strategies, forbearance considerations, and documentation practices.

*Walk me **through** your reasoning step by step – **considering** all **the** relevant factors, trade-offs, **and** bank interests – **before** giving me a final recommendation on how BankUnited should approach a commercial **real** estate client who **is** approaching loan maturity during a period **of** elevated rates **and** declining collateral values.*

What to notice: The response should show its work — market context analysis, borrower position assessment, bank risk exposure, option evaluation — before landing on recommendations. The reasoning chain is the value.

Here are **two** examples **of** how our commercial team has handled similar situations well:

Example **1** – Strong: [paste **a** real **or** hypothetical example **of a** well-handled maturity **extension** negotiation]

Example **2** – Strong: [paste another strong example]

Now give me **a** recommendation **for** how BankUnited should approach **a** commercial real estate client who is approaching loan maturity during **a** period **of** elevated rates **and** declining collateral values. Match **the** style **and** specificity **of the** strong examples.

What to notice: The output should mirror the structure and depth of your examples. If your examples were concise and action-oriented, the new output will be too.

*I need **a** recommendation **on how**
BankUnited should approach
a commercial real estate client who is
approaching loan maturity
during **a** period **of** elevated rates **and**
declining collateral values.
Before you give me **an** answer, ask me **any**
questions you need
to give me **the** most useful possible
recommendation.*

*What to notice: Copilot should ask about
the specific client situation, loan size,
collateral type, borrower relationship
history, and bank appetite for extension vs.
resolution. Your answers will produce a
significantly more tailored
recommendation.*

Play **the** role **of** a skeptical senior credit officer who has seen too many banks extend-**and**-pretend their way **into** larger losses. I'm going **to** propose **that** we offer a **12-month** maturity extension **with** a modified covenant structure **to** a commercial **real** estate client approaching maturity **in** a tough market. Push **back** hard. Tell **me where** this approach **is** risky **and** what I'm **not** thinking **about**.

My proposal: We offer **the** extension **at** current market rate plus a **50bps** risk premium, **with** a requirement **for** a fresh appraisal **and** a debt service coverage covenant **of 1.15x**. If they can't meet DSCR **at** renewal, they agree **to** a principal reduction plan.

What to notice: This output should be genuinely challenging. A good Sparring Partner response will identify risks like appraisal gaming, DSCR covenant

*enforceability, principal reduction
mechanics, and the regulatory optics of the
extension.*

::::

Compare the five outputs. They should be meaningfully different — not just in tone, but in substance, depth, and practical utility. That difference is what you're learning to produce deliberately.

What to Record After the Exercise

Keep notes on:

- *Which technique produced the most immediately useful output for this type of question?*
- *Which technique surprised you most?*
- *Where would you combine techniques? (Role-Based + Chain-of-Thought is a particularly powerful combination for complex analytical tasks.)*
- *What would you add to refine each prompt further?*

Your prompting instincts improve with every iteration. The professionals who become power users of Microsoft 365 Copilot are not the ones who got lucky on the first try — they are the ones who treated each prompt like a hypothesis to be tested and refined.

8. Putting It Together: Your First Week Prompting Plan

This course gives you the frameworks. What moves you from "I understand the techniques" to "I use them automatically" is repetition — applied to real BankUnited work, on real tasks, this week.

Here is a concrete five-day prompting plan:

```
:header-rows: 1

* - Day
  - Technique
  - Task to Try at BankUnited
* - Monday
  - Role-Based Prompting
  - Assign a relevant professional role and
    use it to analyze a document or memo you're
    currently working on
* - Tuesday
  - Chain-of-Thought
  - Use Copilot to reason through a decision
    or recommendation you need to make this week –
    ask for the step-by-step reasoning before the
    conclusion
* - Wednesday
  - Few-Shot
  - Pull two or three examples of a document
```

type your team produces regularly, and ask Copilot to generate a new one in the same style

- * - Thursday
 - Reverse Prompting
 - Let Copilot interview you for a presentation or proposal you need to develop – answer its questions, then review the output
- * - Friday
 - Sparring Partner
 - Present one of your current ideas or recommendations to Copilot in the Sparring Partner mode and see what comes back

By Friday, you will have hands-on experience with all five techniques applied to real work. That is more practical prompting practice than most Copilot users accumulate in their first three months.

9. Productive Struggle Problem

You are a senior relationship manager at BankUnited. Your most important portfolio client — a \$22M commercial real estate borrower with six properties across South Florida — has just

informed you that a competitor bank is offering to refinance their entire portfolio, claiming it can save them \$180,000 annually in interest costs.

You have thirty minutes before the client calls to discuss. You need: a clear understanding of whether the competitor's claim is plausible, a retention strategy that does not simply concede on rate, and talking points that draw on the eight-year relationship history.

You have access to Microsoft 365 Copilot and the five techniques from this chapter.

The challenge: Design a sequence of Copilot prompts — using at least three of the five techniques — that you would actually run in the next twenty-five minutes to prepare for this call. Write out each prompt in full, explain which technique it uses and why you chose it for that step, and describe what you expect from each output.

There is no single right answer. There are better and worse sequences, and the quality of your thinking about *why* you're choosing each technique is as important as the prompt itself.

Discussion

Prompting as a Professional Skill

The five techniques in this chapter are not software features — they are communication frameworks that work because of how language models process context. As Microsoft continues to evolve Microsoft 365 Copilot, the specific mechanics will change. The principles will not.

Consider: In what ways does prompting Copilot resemble briefing a highly capable junior colleague? In what ways does it differ? What does that comparison reveal about where human judgment remains irreplaceable in the AI-assisted workflow?

Discussion Guidelines:

Your response should engage substantively with the comparison between prompting AI and briefing a human colleague. Include at least one specific example from your own professional experience — either an experience where clear communication dramatically improved a colleague's output, or a case where a vague

request produced work that missed the mark entirely. Apply that experience to what you've learned about prompting.

Support your perspective with at least one credible source — this might be something from the Microsoft 365 Copilot official documentation, the Skilling Center resources at adoption.microsoft.com/copilot/skilling-center/, or a relevant piece of research on human-AI collaboration in professional settings.

After posting your response, engage with **at least two classmates** by extending or challenging a specific claim they made — not just affirming it. "I agree because..." is not a sufficient peer response. "I'd push back on your claim that X because I've seen Y, which suggests Z" is.

Do not summarize or repeat what you read.

Share what you think — grounded in evidence — about what this means for how BankUnited professionals should work going forward.

Glossary

Microsoft 365 Copilot

The paid, enterprise-grade AI assistant available **through** BankUnited's Microsoft 365 license. Unlike Copilot Chat, M365 Copilot connects **to** your organizational data via Microsoft Graph **and** operates within your M365 security boundary.

Copilot Chat

The free, general-purpose AI assistant available **at** copilot.microsoft.com. Provides web-connected AI assistance **but does not** access organizational data (emails, files, Teams chats) **by** default.

Microsoft Graph

The intelligence layer within Microsoft 365 **that** maps organizational data – emails, meetings, chats, calendar events, files, **and** shared documents. Copilot uses Microsoft Graph **to** ground **its** responses **in** your actual work context.

Grounding

The process **by** which Microsoft 365 Copilot accesses Microsoft Graph **to** retrieve relevant organizational context **before** generating a response. Grounding **is** what makes Copilot responses personalized **to** your actual work rather than generic.

Copilot Agent

A purpose-built AI assistant configured **for** a specific workflow, knowledge base, **or** organizational function. More specialized than **the** general M365 Copilot experience; configured **through** Microsoft Copilot Studio.

m365.cloud.microsoft

The unified web portal **for** Microsoft 365 Copilot – **the "front door" that** provides access **to** all M365 Copilot capabilities, **the** Prompt Gallery, **and** connected M365 apps **from** a single authenticated session.

Microsoft 365 Copilot Search

A universal search capability within **the** M365 Copilot experience **that** searches across all connected M365 apps (Outlook, SharePoint, Teams, OneDrive) **and** **third-party** data **in** a single query.

Role-Based Prompting

A prompting technique **that** assigns a specific professional identity **or** perspective **to** Copilot **before** making a request, shifting **the** analytical frame, vocabulary, **and** focus **of** **the** output.

Chain-**of**-Thought Reasoning

A prompting technique **that** asks Copilot **to** show **its** reasoning step-**by**-step **before** giving a final answer, producing more transparent, coherent, **and** auditable analytical output.

Few-Shot Prompting

A prompting technique **that** provides Copilot **with** a small **set of** high-quality examples **before** asking **it to** produce new output – teaching **by** demonstration rather than **by** description.

Reverse Prompting

A prompting technique **that** inverts **the** standard dynamic **by** asking Copilot **to** interview you **with** clarifying questions **before** producing output, ensuring **the** request **is** well-understood **before** work begins.

Sparring Partner

A prompting technique **that** assigns Copilot **the** role **of** a skeptical, critical counterpart who pushes **back on** your ideas – used **to** pressure-test thinking **before** high-stakes conversations.

Goldilocks Zone

The optimal prompt **length and** specificity **for** a **given** task – precise enough **to** produce relevant output, concise enough **to** avoid confusing **or** overloading **the** model. Typically achieved **with the** four-component prompt structure: Context, Goal, Format, Constraints.

Four-Component Prompt

A prompting structure **with** four elements – Context, Goal, Format, **and** Constraints – **that** consistently produces more useful output than

vague **or** overly complex prompts.

Prompt Gallery

Microsoft's official curated library **of** tested, verified prompts organized **by** job function **and application**. Available **at** m365.cloud.microsoft/copilot-prompts.

Copilot Skilling Center

Microsoft's official learning hub **for** Microsoft 365 Copilot, **with** structured learning paths **and** role-specific guidance. Available **at** adoption.microsoft.com/copilot/skilling-center/.

Permission Scoping

The architectural principle **that** Microsoft 365 Copilot only surfaces information **the** signed-in user already has permission **to** access – ensuring Copilot cannot bypass existing security **and** compliance controls.

Leader's Takeaway

Microsoft 365 Copilot is not a feature you turn on. It is a capability you develop — in yourself, and in your team.

The infrastructure is in place. The security architecture is sound. The intelligence layer that grounds Copilot in your organizational context is active. The tools are available at m365.cloud.microsoft right now.

What remains is the skill — specifically, the prompting skill that determines whether Copilot becomes a transformative part of how BankUnited works or sits as an underutilized line item in the software budget.

The five techniques in this chapter are not exotic. They are communication principles applied to a new kind of tool. Role-Based Prompting is how you brief a specialist. Chain-of-Thought is how you ask for a memo, not just a verdict. Few-Shot is how you onboard a new analyst with examples. Reverse Prompting is how you let the expert interview you before the work begins. The Sparring Partner is how you moot-court your best thinking before the real conversation.

These techniques are learnable. They improve with repetition. And the professionals on your team who develop fluency with them fastest will

not just be more productive individually — they will become the informal teachers who spread that capability to everyone around them.

That is the compound return on investing in this skill. It is not linear — it multiplies.